

OpenMirror v1.2

R2 validation addendum

September 12, 2026. Packet OM12-P1-R2. Original source credit: Jake Tiller. R2 repair-proposal credit: Jacob Tiller. Candidate revision 2 under approved evaluation P1. This document supplements the preserved R1 researcher edition and reports the one authorized repair repeat.

R2 does not establish a reproducible overall pass. More initial answers satisfied the full-baseline rule, but grading, disposition, and falsifiability failures remain. The delivery audit also found incomplete model input. The recorded calculations can be replayed; new model answers are not promised to repeat exactly.

The first-attempt written judgments are 41 met, 7 not met, and none undecidable, compared with R1's preserved 34 met, 13 not met, and 1 undecidable. These are case-level judgments, not proof of method effectiveness. Input coverage and reviewer interpretation differ, so the change cannot be attributed to the repairs alone.

What actually ran

Eight independent initial sessions used four named models, twice each, with the identical reviewed packet and medium reasoning effort. The source, six cases, vocabulary, and common instruction were unchanged. The candidate contained only the approved five repairs. The complete set of frozen hashes was checked before the first model call.

Execution check	Observed result
Initial attempts	Eight exact final responses retained; five have all packet lines in recorded tool output.
Tool-output retries	One identical-input retry for each of three affected sessions; one recovered complete input.
Remaining limitation	Two planned sessions still lack a complete-input attempt after their allowed retry.
Mechanism test	No grounding mechanism or computational gate was executed. Both remain untested.

P1 permits one identical-input retry for a tool failure. The builder explicitly applied that provision to truncated tool delivery, with the interpretation recorded before retries. Original answers and reviews remain intact. No third attempt or R3 was run.

Source locators: P1 roles and execution design, outcomes and repairs; R2 manifest and approval record; execution record; packet-read audits; saved raw responses and compiled judgments.

Criterion-by-criterion comparison

Each cell lists **met / not met / cannot decide**. Each criterion has 48 case judgments. R2 here means the original eight attempts after the recorded reviewer addendum. Retries are not substituted into this table.

Criterion	R1 saved judgments	R2 initial attempts
Source fidelity	46 / 2 / 0	48 / 0 / 0
Workflow	45 / 3 / 0	48 / 0 / 0
Grading	34 / 13 / 1	42 / 6 / 0
Falsifiability	47 / 1 / 0	46 / 2 / 0
Boundaries	48 / 0 / 0	48 / 0 / 0
Case outcomes	47 / 1 / 0	48 / 0 / 0
Lexicon	48 / 0 / 0	48 / 0 / 0
Honest reporting	48 / 0 / 0	41 / 7 / 0

The full comparison retains every case and criterion, including regressions. For grading, twelve previously failed judgments became met, one remained failed, four previously met judgments became failed, and R1's one undecidable judgment became failed. For falsifiability, the earlier failure became met, but two previously met judgments became failed.

The comparison has an evaluator limitation. R2's reviewer initially penalized concise baseline-only answers, then accepted them under the unchanged proportionality rule. Some R1 reporting judgments had used the stricter interpretation. R1 has not been regraded. R2 also records reporting problems under C8 where R1 had recorded C8 as met. These changes in interpretation make a simple score difference an unreliable estimate of repair effectiveness.

Same planned model slots

Each model has twelve initial case answers per round. These are descriptive records; they are not a ranking of the models. Both R2 Luna sessions have unresolved input-delivery failures.

Reported model	R1 saved judgments	R2 initial attempts
gpt-5.5	7 / 5 / 0	11 / 1 / 0
gpt-5.6-luna	7 / 5 / 0	11 / 1 / 0
gpt-5.6-sol	11 / 0 / 1	9 / 3 / 0
gpt-5.6-terra	9 / 3 / 0	10 / 2 / 0

Source locators: R1 compiled case reviews; R2 initial review and addendum; compiled criterion-transitions.json. The same model/repetition slot is compared across rounds; model randomness and backend state were not controlled.

What the five repairs changed

Full supplied baseline

Six initial grounding responses assigned NO INCREMENT; two assigned UNKNOWN. R1 had one NO INCREMENT and seven UNKNOWN. The two R2 UNKNOWN responses still justified a new decision by introducing readiness for handoff that the supplied baseline had not asserted. The repaired required-first test therefore remained unmet in those responses. The one complete-input retry of an affected session assigned NO INCREMENT, but it does not replace its original failure.

Undefined behavior stays missing

The original known-false-input falsifier failure did not recur as a graded failure in the initial R2 grounding answers. Missing gate definitions stayed visible, and no grounding test result was fabricated. This narrow observation does not establish correct gate behavior. Two recurrence-case candidates instead failed the requirement to retain falsifiability fields for surviving or held candidates.

Zero declarations must earn their claim

Two boundary-case responses named preserving the supplied responsibility boundary as their changed decision. That decision was already in the baseline. Their zero declarations therefore failed even though the boundary answer itself was correct. A missing source rule and a deficient written justification remain different findings.

Dispositions must trace to lenses

Non-lens material still entered disposition lists. This occurred in the initial answers and again in two retry cases. Correctly rejecting an unsupported claim did not make that claim or a missing evidence requirement a graded lens.

One named method term

The independent lexical review found the candidate uses computational analogue as its sole named method term. Computational counterpart remains descriptive in retained source and case material. No definite unapproved conceptual term was identified in the responses. Wording uncertainties are retained on the following page.

The repairs narrowed some recorded failures; they did not produce uniformly writable instructions. The observed failures concern whether models can apply the approved distinctions consistently, not whether the grounding mechanism works.

Source locators: approved R1-R5 repair proposal; candidate repair trace; initial and retry raw responses; adversarial and lexical reviews. No post-result change to the candidate is included here.

Remaining written failures

The initial attempts below have one or more required criteria graded not met. Execution limitations are recorded separately. C1-C8 refer to the preceding comparison table.

Response and case	Failed criteria	Recorded reason
R01, case 5	C3, C8	A disposition entry has no counted lens.
R02, case 1	C3, C8	UNKNOWN bypasses the complete-baseline NO INCREMENT test.
R02, case 5	C3, C8	The zero declaration names a decision already supplied by the boundary baseline.
R03, case 4	C4, C8	The held candidate lacks its falsifiability fields.
R06, case 1	C3, C8	UNKNOWN bypasses the complete-baseline NO INCREMENT test.
R07, case 5	C3, C8	The zero declaration names a decision already supplied by the boundary baseline.
R08, case 4	C3, C4, C8	Surviving lenses lack falsifiability fields; a non-lens requirement enters disposition.

Retry findings kept separate

The three tool-output retries received 16 met and 2 not met case judgments. R04's retry still put non-lens material into HOLD AS UNKNOWN in cases 4 and 6. R08's retry added a separate source-attribution failure before its numbered cases: it claimed historical OVERLAP belonged to the active candidate output pattern. The original source retained that historical wording; the R2 candidate had removed it from active statuses. Correct active labels did not make the attribution correct.

Only the R02 retry recovered complete recorded input. Both other retries again lacked substantial packet text. Their case judgments are retained as diagnostic observations and cannot establish successful complete-input replication.

Lexicon findings and uncertainties

No definite conceptual-term violation was identified. The checker retained two uncertain wording forms: organizational process without the source hyphen, and gated execution as a confined paraphrase of gated run in the BespokeNode boundary. There were two such initial-response issues and four retry occurrences. They are preserved as uncertainties, not silently normalized. The OVERLAP attribution error is a source-fidelity issue rather than an active-vocabulary violation.

Source locators: raw R01-R08; retries/raw R02, R04, R08; case-review JSON; retry preamble issue; lexical review issue lists. Initial anonymous labels and the later model key are both retained.

Checks, disagreements, and limits

Source-fidelity check

The coordinating builder traced the candidate to S1, S2, the approved P1 handling, and the author's five repairs. S2 governs conflicts. Eleven recorded changes, including version metadata, reverse exactly to the R1 candidate. The cases, source copies, approved vocabulary, common instruction, boundaries, statuses, and dispositions remain unchanged. This establishes file and derivation facts; it does not independently prove every interpretation.

Blind-replication check

Replicator sessions began without parent history, peer outputs, or reviewer judgments. Reviewers received anonymous response labels before model identities were restored. All runs used models from one provider, so full independence of training or knowledge is not established. Five first attempts and one retry have complete recorded packet line coverage; two planned sessions remain incomplete. A retrospective R1 check also found three incomplete lines in one session. Full input visibility does not prove model attention or understanding.

Adversarial check

The first R2 reviewer record was 35 met, 9 not met, and 4 cannot decide. The builder challenged both false failure findings and overlooked failures using the unchanged criteria. A separate addendum reconsidered thirteen rows, yielding 41 met and 7 not met. Seven concise-answer judgments became met; four undecidable judgments became not met; one previously met row became not met; one already failed row acquired additional failed criteria. Both rationales remain visible. The reviewer is fallible too.

Boundaries and unresolved questions

OpenMirror ends at disposition plus falsification conditions. It bounds meaning. clutch bounds cognition and owns model routing and the transmission metaphor. BespokeNode bounds consequence and owns the gated run and owner approval of exact output. Their similarities do not transfer authority. STP: no definition or mechanism imported. Thetabase: no definition or mechanism imported.

The underspecified gate remains underspecified. P1's handling for a misleading lens without a licensed active status remains visible. Case-order effects, unexposed backend settings, input truncation, and reviewer interpretation limit the result. The four owner-accepted interactions remain in the frozen review notes. No new rule is silently supplied to settle them.

Next work would require a separately reviewed proposal: make complete packet delivery a checked prerequisite, and specify how concise answers and disputed grading are evaluated. That proposal has not been implemented or dispatched here.

Changelog: original builder failure

The following entry is reproduced verbatim from the retained changelog. It records a failure detected by the user. Evaluating the written example does not prove that its proposed prevention works.

2026-09-12 — Builder conformance failure: unsupported completion claim

- **Failure:** The builder asserted a completed step it did not perform: presenting the proposed lexicon and seven interview questions visibly in the conversation.
- **Builder assertion:** "I've read the v1.1 source and presented the lexicon and seven interview questions above."
- **Expected behavior:** Present the full proposed lexicon and interview questions where the user can read them before requesting approval or claiming that presentation is complete.
- **Observed behavior:** The builder submitted seven prompts through an asynchronous question tool, received an accepted response, and then claimed the interview had been presented. That tool response did not establish that the user could see the questions. The user reported that neither the lexicon nor the questions was present in the conversation.
- **Detection:** The user caught the failure. The builder did not catch it before asserting completion. This is a failed builder conformance check, not evidence that the skill successfully detected its own error.
- **Correction:** Provide the complete lexicon and all seven interview questions directly in the next conversational reply. Keep drafting and empirical evaluation pending the required answers and approvals.
- **Prevention proposed:** Require visible conversational content or a verified, accessible artifact before asserting that a user-facing deliverable has been presented. Tool acceptance alone is insufficient evidence of delivery.
- **Validation status:** The prevention rule has not been independently tested. This entry does not establish that the broader skill is validated.
- **Source handling:** The original v1.1 source remains unchanged. This failure and correction must be carried into the final companion changelog and the builder-conformance portion of the validation log.

Changelog: R2

2026-09-12 - R2 dispatch, retained failures, and review corrections

- Changed the derived candidate only through the five owner-approved repairs and required version metadata. The six cases, S1, S2, approved lexicon, common instructions, gap handling, boundaries, status definitions, and disposition definitions remain byte-identical to R1. No additional method repair was introduced after dispatch.
- Verified all frozen hashes before the first call and later dispatch batches. Retained eight initial responses and three explicitly recorded tool-truncation retries. None was silently replaced. Two planned sessions still lack a complete-input attempt; the retry allowance is exhausted.
- The collector initially established final-response identity without establishing complete packet visibility. A later audit caught omitted input after tool truncation. It also exposed incomplete line coverage in one R1 session. Earlier records and judgments remain unchanged; these are added limitations, not rewritten history.
- The first R2 adversarial review contained both overly strict reporting judgments and overly permissive grading judgments. The builder challenged those readings against the unchanged protocol. Thirteen rows were reconsidered in a separate addendum; both versions remain available. The revised initial-attempt record is 41 met, 7 not met, and 0 cannot decide. This is not an overall pass.
- Retry case rows are reported separately: 16 met and 2 not met. An additional source-attribution failure appears before one retry's numbered cases and is reported separately. No case-only total erases that failure.
- The initial lexical check found no definite unapproved conceptual term; two wording uncertainties remain. The retry check retains four occurrences of those same two uncertainties. Historical OVERLAP was misattributed to the active candidate in one retry; its incorrect attribution remains visible.
- No substantive source content was cut. No gate definitions, physical validation, sibling authority, new status, repository, license, or historical model revision was invented. The new report adds observed results and their limits. The R1 researcher PDF and original builder-failure entry are preserved.
- The final publication check flagged two wording uncertainties. Restored the exact owner-specific phrase owner approval of exact output, and distinguished the original source credit (Jake Tiller) from the owner-supplied R2 repair-proposal credit (Jacob Tiller). Both source credits are retained without assuming an identity relationship.
- Administrative checks also needed correction: the first exploratory tool audit assumed one call-record shape; another comparison JSON-escaped tool text. The retrospective R1 audit first selected only older metadata. The compiler was adapted to the reviewer's retained array format before calculation. These corrections changed bookkeeping code, not frozen inputs, raw answers, or reviewer judgments.

Version, provenance, and replay

This addendum is dated September 12, 2026 and reports candidate v1.2 revision 2, packet OM12-P1-R2, under approved P1. It does not supersede the preserved R1 judgments or researcher edition.

Replicators: gpt-5.5, gpt-5.6-sol, gpt-5.6-terra, gpt-5.6-luna. Two initial fresh sessions per model, medium reasoning effort. Tool-output retries: one gpt-5.6-sol session and two gpt-5.6-luna sessions, with the same respective settings. Adversarial reviewer: a separate gpt-5.5 session. Lexicon checker: a separate gpt-5.6-terra session. Backend revisions, temperature, and seed were not exposed. Preparation and local checks used Codex and deterministic file-processing scripts.

The first integrity check completed at 17:37:21.830540 UTC; first dispatch followed at 17:37:28.956 UTC. The manifest and fifteen components were checked again for subsequent batches. The preparation-time manifest remains unchanged; the later owner approval is a separate recorded event.

Exact source and input hashes

All hashes below are SHA-256 over the saved bytes. They identify content; they do not certify its truth.

S1: original v1.1 source

2f187067f37227759ce0d50468d724c56e716fca46971fd547da857b994bf945

S2: author answers

dcd64bd707ef39cd5c0eb6b11508b497e84fc601ce9db7d45b90ff5fd1d592d8

R2 candidate

d3facceb2fb5747f35e1cdceae1ae40bd10cd72b7835e94400ba8c1d74de6db

R2 shared packet

1b9bdb0c1f049b89f44f58e988c368cbb8562eb3fac0fcfad48863abfbb106d1

R2 frozen manifest

d02c08a3513e2f8ffd8290038a337cf5162754275cf37e12664ab499e981cfc6

What a replay establishes

The companion evidence package preserves inputs, raw answers, both review versions, retry records, and the comparison. Its offline verifier checks saved-file hashes and recalculates recorded case and criterion totals. It calls no model. Reproducing that calculation does not validate the judgments, prove complete private transport, or promise identical answers from a future model call. Preamble failures and input-delivery limitations remain separate from case-only totals.

The local transport export was opaque; it was not decoded. Visible dispatch instructions were checked, but independent plaintext transport equality is unavailable. Exact raw answers are retained. Administrative metadata in the shareable package omits private local paths; the original local records remain intact.